

Retrieval Pipelines

A retrieval pipeline combines several lanes: a document router, a parent summary lane, and a global child lane. Document routing is never a recall gate: a child hit survives even when its document scores zero.

The pipeline fuses per-lane rankings with reciprocal-rank fusion. Each hit carries the representation kind and contract that produced it, so provenance is auditable at every stage.

A retrieval pipeline must descend deterministically from a routed document to its parent and child evidence. Document identity, parent identity, and chunk identity are linked by stable identifiers rather than filename guessing.

The pipeline serves answers through evidence bundles. Graph evidence augments textual retrieval; textual passages answer independently when the graph has nothing to contribute.